

DATA MANAGEMENT SYSTEMS

PROJECT 2 ANALYSIS OF A REAL WORLD PROBLEM

USED CAR SALES IN THE USA AND RUSSIA

Amelia Waszkowska

1. Objective

The goal of this project is to develop a machine learning model using big data analysis techniques to predict the prices of used cars based on various features and attributes. By leveraging open datasets related to used car offers, the project aims to explore patterns, correlations, and factors influencing car prices, ultimately enabling accurate price predictions.

2. Datasets

I used following big datasets available on kaggle:

<https://www.kaggle.com/datasets/ekibee/car-sales-information>

<https://www.kaggle.com/datasets/rupeshraundal/marketcheck-automotive-data-us-canada?select=us-dealers-used.csv>

3. Storage

Massive data require using specific data processing framework, like Apache Hadoop and or Apache Spark. These two frameworks are often used together:

Hadoop HDFS (Hadoop Distributed File System) is a distributed file system designed to store and manage large volumes of data across multiple computers in a Hadoop cluster. It is a key component of the Apache Hadoop framework, which is widely used for big data processing and analytics.

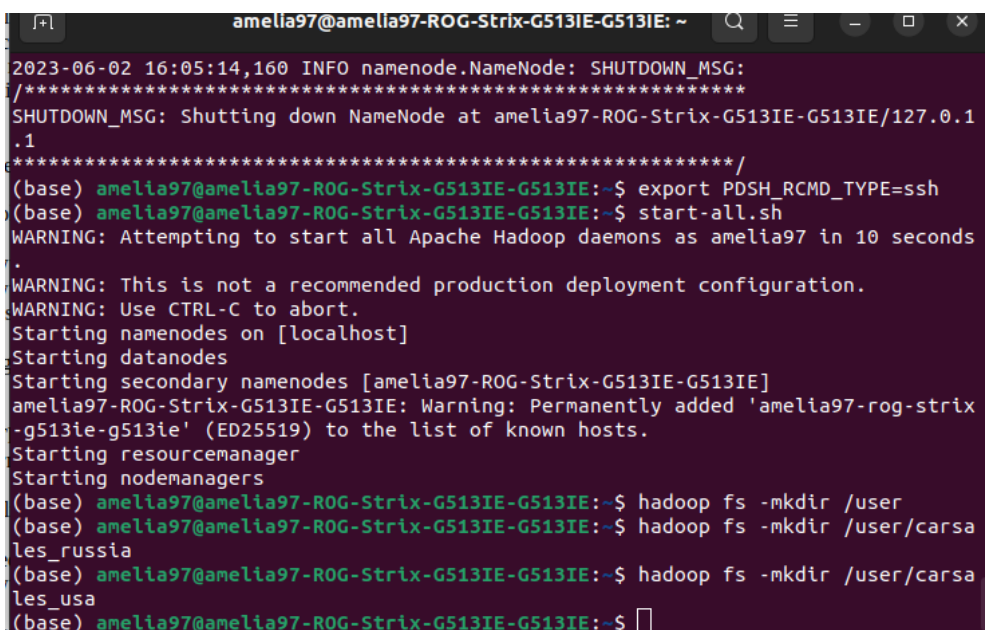
Apache Spark is an open-source distributed computing framework that is designed for big data processing and analytics. It provides an interface for programming and executing data processing tasks across a cluster of machines, making it highly scalable and efficient.

4. Installing Hadoop

I installed Hadoop 3.2.4 following this tutorial:

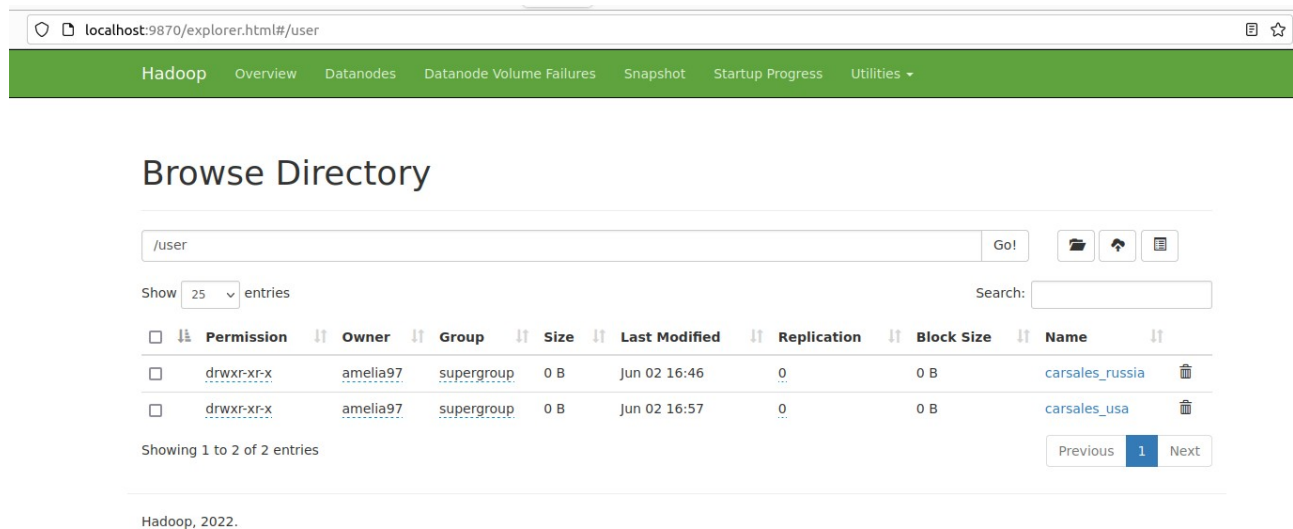
<https://www.youtube.com/watch?v=Slbi-uzPtnw>

Next, I created two distinct directories in my Hadoop HDFS for storing my csv files:

A terminal window screenshot showing the installation and configuration of Hadoop. The terminal is titled 'amelia97@amelia97-ROG-Strix-G513IE-G513IE: ~'. It shows a shutdown message for the NameNode, followed by the user running 'export PDSH_RCMD_TYPE=ssh' and 'start-all.sh'. A warning message indicates that all Apache Hadoop daemons are being started as 'amelia97' in 10 seconds. Another warning states that this is not a recommended production deployment configuration and that the user should use CTRL-C to abort. The terminal then shows the starting of namenodes on [localhost], datanodes, secondary namenodes [amelia97-ROG-Strix-G513IE-G513IE], and resource manager. Finally, the user runs 'hadoop fs -mkdir /user/carsales_usa' and 'hadoop fs -mkdir /user/carsales_russia' to create the directories for storing CSV files.

```
[~] amelia97@amelia97-ROG-Strix-G513IE-G513IE: ~
2023-06-02 16:05:14,160 INFO namenode.NameNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down NameNode at amelia97-ROG-Strix-G513IE-G513IE/127.0.1
.1
*****/
(base) amelia97@amelia97-ROG-Strix-G513IE-G513IE:~$ export PDSH_RCMD_TYPE=ssh
(base) amelia97@amelia97-ROG-Strix-G513IE-G513IE:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as amelia97 in 10 seconds
.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [amelia97-ROG-Strix-G513IE-G513IE]
amelia97-ROG-Strix-G513IE-G513IE: Warning: Permanently added 'amelia97-rog-strix-g513ie-g513ie' (ED25519) to the list of known hosts.
Starting resource manager
Starting nodemanagers
(base) amelia97@amelia97-ROG-Strix-G513IE-G513IE:~$ hadoop fs -mkdir /user/carsales_usa
(base) amelia97@amelia97-ROG-Strix-G513IE-G513IE:~$ hadoop fs -mkdir /user/carsales_russia
(base) amelia97@amelia97-ROG-Strix-G513IE-G513IE:~$
```

After opening the localhost 9870 in the browser, it looks like this:



Next, I uploaded the files:

```
les_usa
(base) amelia97@amelia97-ROG-Strix-G513IE-G513IE:~$ hadoop fs -put /home/amelia97/Downloads/region25_en.csv /user/carsales_russia
(base) amelia97@amelia97-ROG-Strix-G513IE-G513IE:~$ hadoop fs -put /home/amelia97/Downloads/region41_en.csv /user/carsales_russia
(base) amelia97@amelia97-ROG-Strix-G513IE-G513IE:~$ hadoop fs -put /home/amelia97/Downloads/us-dealers-used.csv /user/carsales_usa
(base) amelia97@amelia97-ROG-Strix-G513IE-G513IE:~$
```

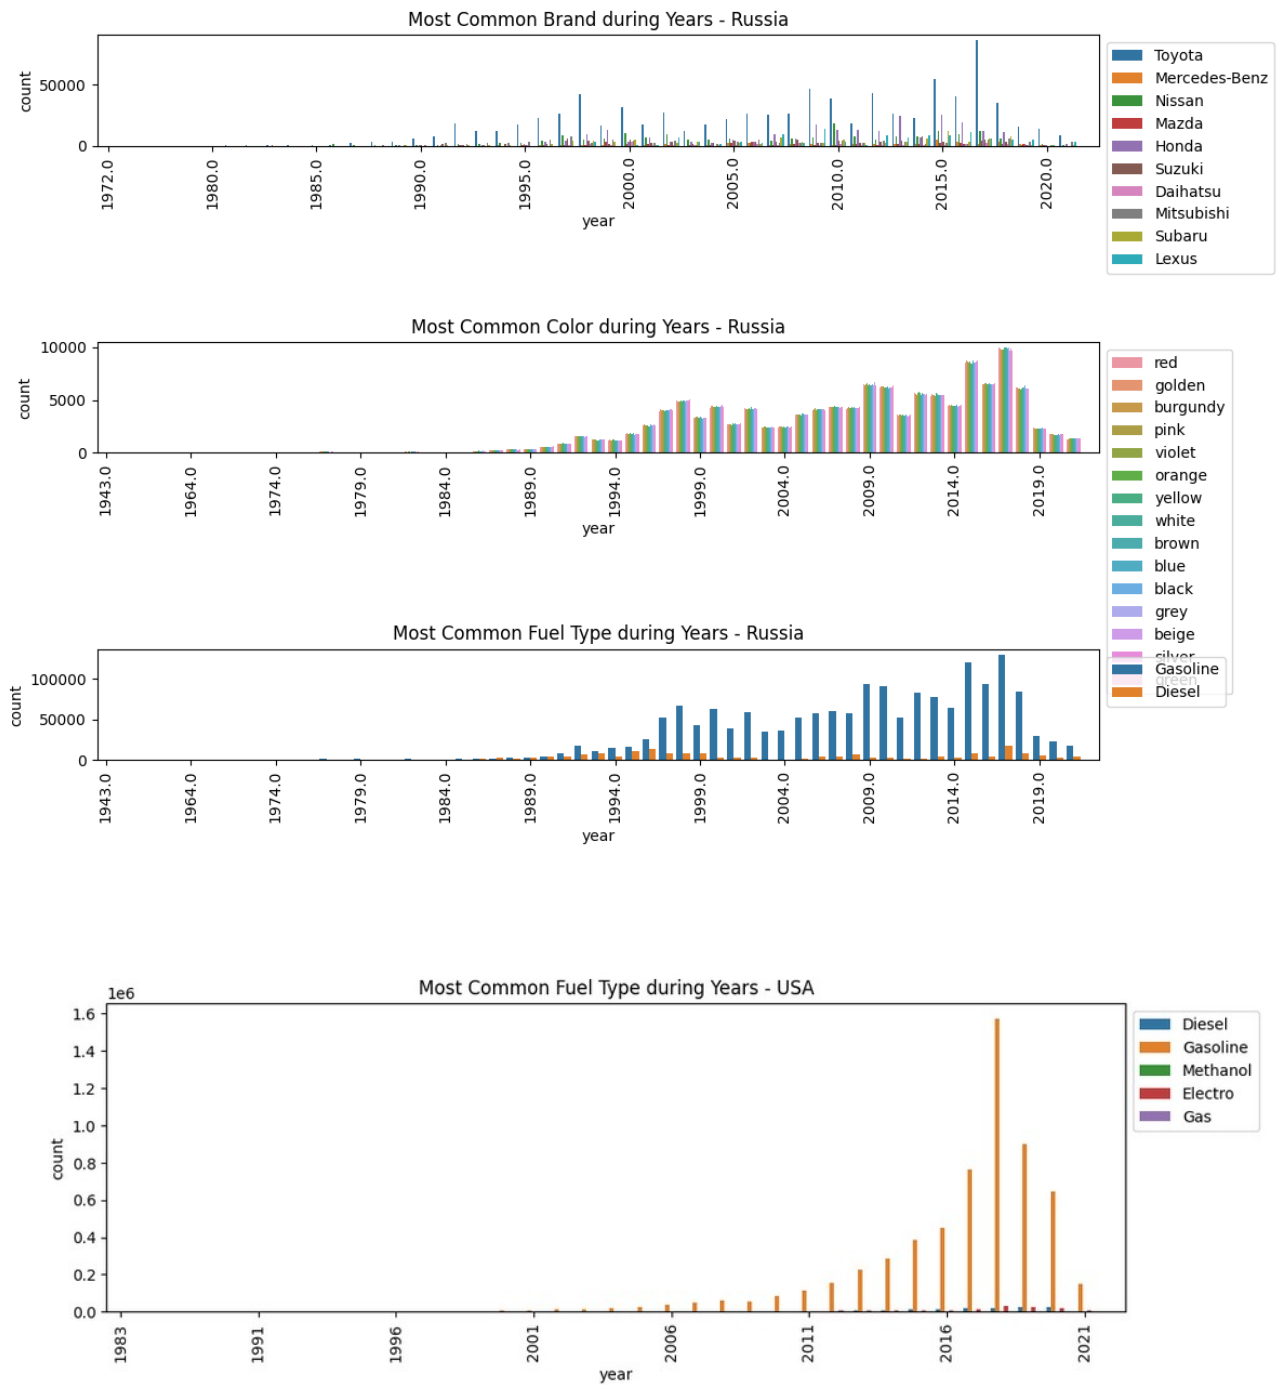
5. Setting up Spark and Jupyter


Since Spark may be problematic to install, it is advisable to create a new environment for Jupyter notebook, in which we want to use Spark. So I did. To use Spark in Jupyter notebook, we can just run "!pip install" in the notebook. Of course, this installation approach has a limitation: it only installs pyspark within the notebook environment, and other Jupyter Notebook sessions or terminals won't have access to it. But in this project it's not a problem.

From now and on, I will write majority of descriptions in the Jupyter notebook.


6. Pipeline decription

7. Most common brand, color and fuel type throughout the years.







8. Median car price




Median Car Price by Year - Russia





9. Average price for fuel type



10. Average price for brand



11. Check if data follow a normal distribution (examples)



Generally, the dataset is skewed and doesn't follow the normal distribution.

12. The lack of some USA plots and problem with running out of the memory.

During the project I encountered the following error:

```
23/06/04 20:03:43 ERROR Executor: Exception in task 2.0 in stage 46.0 (TID 360)]
java.lang.OutOfMemoryError: GC overhead limit exceeded
    at org.apache.spark.unsafe.types.UTF8String.fromAddress(UTF8String.java:132)
    at org.apache.spark.sql.catalyst.expressions.UnsafeRow.getUTF8String(UnsafeRow.java:406)
    at org.apache.spark.sql.catalyst.expressions.GeneratedClass$SpecificOrdering.compare_0_0$(Unknown Source)
    at org.apache.spark.sql.catalyst.expressions.GeneratedClass$SpecificOrdering.compare(Unknown Source)
    at org.apache.spark.sql.execution.UnsafeKVExternalSorter$KVComparator.compare(UnsafeKVExternalSorter.java:27
3)
    at org.apache.spark.util.collection.unsafe.sort.UnsafeSorterSpillMerger.lambda$new$0(UnsafeSorterSpillMerge
r.java:37)
    at org.apache.spark.util.collection.unsafe.sort.UnsafeSorterSpillMerger$$Lambda$4122/991538801.compare(Unkno
wn Source)
    at java.util.PriorityQueue.siftDownUsingComparator(PriorityQueue.java:719)
    at java.util.PriorityQueue.siftDown(PriorityQueue.java:687)
    at java.util.PriorityQueue.poll(PriorityQueue.java:595)
    at java.util.AbstractQueue.remove(AbstractQueue.java:113)
    at org.apache.spark.util.collection.unsafe.sort.UnsafeSorterSpillMerger$1.loadNext(UnsafeSorterSpillMerger.j
ava:91)
    at org.apache.spark.sql.execution.UnsafeKVExternalSorter$KVSorterIterator.next(UnsafeKVExternalSorter.java:2
8)
```

This means that there is not enough memory to process the data. It happened several times during processing of the USA dataset, which is significantly bigger.

It's known that Spark is memory-centric. That out of memory exception can occur at the Driver or Executor level. The Driver is a Java process, where the main() runs. It submits jobs, converts jobs to stages and coordinates tasks between executors, creates a SparkSession etc. Executors are launched at the start of a Spark Application with the help of cluster Manager. It runs an individual task and returns the result to the Driver. It can also persist data in the worker nodes for re-usability.

This kind of error can occur due to incorrect usage of Spark. The driver is only an orchestrator/coordinator, the executors are provided with more memory to operate. We can even set a proper limit using spark.driver.maxResultSize.

Out of memory at the Executor level on the other hand, occurs usually because of either inefficient queries, high concurrency or incorrect configuration.

In case of high concurrency, there will be inappropriate number of spark cores for our executors, and we will have to process too many partitions. All these will be running in parallel and will have their memory overhead; therefore, they need the executor memory and can probably cause OutOfMemory errors. To fix this, we can configure spark.default.parallelism and spark.executor.cores and based on your requirement you can decide the numbers.

In case of incorrect configuration, there is a possibility that the application fails due to YARN memory overhead issue (if Spark is running on YARN). Therefore, based on each requirement, the configuration has to be done properly so that output does not spill on the disk. Configuring memory using spark.yarn.executor.memoryOverhead will help to resolve this.

Since I'm configuring Spark application like this:

```
import os
memory = '20g'
pyspark_submit_args = '--driver-memory ' + memory + ' pyspark-shell'
pyspark_submit_args = '--executor-memory ' + memory + ' pyspark-shell'
os.environ["PYSPARK_SUBMIT_ARGS"] = pyspark_submit_args
```

```
spark = SparkSession.builder \
    .config("spark.executor.instances", "4") \
    .appName("CarSales") \
```

.getOrCreate()

In my case the reason of the error is that Spark in Jupyter deos not read those configurations properly, as we can see going to the Spark UI after configuration of the app:

The screenshot shows the Spark UI 'Executors' tab. At the top, there's a navigation bar with 'Jobs', 'Stages', 'Storage', 'Environment', and 'Executors'. Below it, a 'Summary' section contains a table with 14 columns: RDD Blocks, Storage Memory, Disk Used, Cores, Active Tasks, Failed Tasks, Complete Tasks, Total Tasks, Task Time (GC Time), Input, Shuffle Read, Shuffle Write, and Excluded. The table has three rows: 'Active(1)', 'Dead(0)', and 'Total(1)'. Below the summary is a 'Show Additional Metrics' link. The main section is titled 'Executors' and has a 'Show 20 entries' dropdown and a 'Search:' input. It contains a detailed table with 16 columns: Executor ID, Address, Status, RDD Blocks, Storage Memory, Disk Used, Cores, Active Tasks, Failed Tasks, Complete Tasks, Total Tasks, Task Time (GC Time), Input, Shuffle Read, Shuffle Write, and Thread Dump. There is one entry for the 'driver' with status 'Active'. At the bottom, it says 'Showing 1 to 1 of 1 entries' and has 'Previous', '1', and 'Next' navigation links.

	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Excluded
Active(1)	0	0.0 B / 366.3 MiB	0.0 B	16	0	0	0	0	26 s (0.1 s)	0.0 B	0.0 B	0.0 B	0
Dead(0)	0	0.0 B / 0.0 B	0.0 B	0	0	0	0	0	0.0 ms (0.0 ms)	0.0 B	0.0 B	0.0 B	0
Total(1)	0	0.0 B / 366.3 MiB	0.0 B	16	0	0	0	0	26 s (0.1 s)	0.0 B	0.0 B	0.0 B	0

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write	Thread Dump
driver	192.168.0.237:45951	Active	0	0.0 B / 366.3 MiB	0.0 B	16	0	0	0	0	26 s (0.1 s)	0.0 B	0.0 B	0.0 B	Thread Dump

It is clear that I have only the driver, no executors. The steps I would take to solve this issue would start with installing Spark from the command line, and configuring it carefully, not like in this case I just !pip installed it in the Jupyter Notebook.

13. Results of the predictions of the Linear Regression model

MAE for rus model: 4663.33656174515
RMSE for rus model: 9984.815625262503

The predictions given by the model are far from satisfactory. Steps that could be taken in order to obtain better results would consist of: normalizing numerical data, reemoving poorly represented areas of features (just as I have done in case of the year column- removing representation under 2010). From this approach however, we can deduce 2 conclusions. First, it is hard to predict prices on the market so diverse as it is in Russia- world's biggest country. Second, it is rprobably impossible to obtain a good model, from such a diversified dataset, containing underrepresented areas. It would most likely possible however, to narrow this dataset to e.g. years 2010-2015, brands to 5 top selling etc., and then obtain a good-performing model.